

Llama Stack

A streamlined developer experience enabling seamless AI application development.

Build Once, Deploy Anywhere

Cookbooks



Docs

What is Llama Stack?

Llama Stack is a comprehensive system that provides a uniform set of tools for building, scaling, and deploying generative AI applications, enabling developers to create, integrate, and orchestrate multiple AI services and capabilities into an adaptable setup.

Llama Stack codifies best practices across the Llama ecosystem. With the release of the Llama 4 herd of models, you can use Llama Stack to build innovative, personalized applications with the most accessible and scalable generation of Llama.

[Overview](#)[Docs](#)[Experience](#)[Llama Protections](#)[Resources](#)

Overview

Llama Stack Provides

- ✓ Flexible Deployment Options
- ✓ Consistent Application Behavior
- ✓ Robust Distribution Partner Network



Standardized APIs

Provides consistent interfaces for building and deploying AI applications.



Flexible Deployment Options

Supports local development, cloud, on-premises, and mobile environments.



Pre-built Tools

Offers tools for inference, safety monitoring, and memory management.



Scalable Infrastructure

Facilitates easy scaling of AI applications.



Strong Partner Network

Collaborates with various providers to offer specialized services.



Telemetry and Monitoring

Includes built-in support for tracing requests, evaluating model outputs.

Our approach



Docs

Need help getting started building on Llama Stack? Check out our documentation.

[> Learn more](#)



Cookbooks

The tools and resources you need to build with Llama Stack. Access prepackaged, verified provider distributions and mobile SDKs with iOS and Android support.

[Learn more](#)

Llama Stack: a streamlined developer experience

Build faster, deploy anywhere and get the most out of the latest Llama models on day 1.

› [Learn more](#)

Developers

Customize models, build agentic apps, etc



Llama Stack

API and CLI



Our partner ecosystem





Llama Protections

Making safety tools accessible to everyone.

Enabling developers, advancing safety, and building an open ecosystem.



[Learn more](#)

Llama Stack resources

INTRODUCTION TO Llama Stack

Introduction to Llama Stack

Learn about Llama Stack from Dalton Flanagan Meta's Software engineering working on AI.

[> Learn more](#)

BUILDING RAG AGENTS IN Llama Stack

Building Agents with Llama Stack

This technical tutorial will show you how to build a RAG agent using Llama Stack.

[> Learn more](#)

Stay up-to-date

Our latest updates delivered to your inbox

Subscribe to our newsletter to keep up with the latest Llama updates, releases and more.





Documentation



Overview

Models

Getting the Models

Running Llama

How-To Guides

Integration Guides

Community Support

Community



Resources



Llama Protections

